

# INTRODUCTION TO DEEP REINFORCEMENT LEARNING

## IA FRAMEWORKS

---

# TOOLS

## ML Python Libraries



## Python Environment



## Viz' Python Libraries



## Framework & Tool



# TABLE OF CONTENTS

---

Introduction

Definitions

Policy Gradient algorithm

# INTRODUCTION


---

# DEFINITIONS

In **reinforcement learning**, an agent makes **OBSERVATIONS** of the **STATES** of an **ENVIRONMENT**.

It takes **ACTION** within the **ENVIRONMENT** and receives **REWARDS**.

## Example: Walking Robot



# DEFINITIONS

When using reinforcement learning you need to define:

An **AGENT**. The one who will take **action** within the **environment**.

An **ENVIRONMENT** where the **action** will be taken.

The **STATE** of the **environment** defines it after each **action**.

The **OBSERVATION** is what we will use from the **state** to decide the next **action**.

A list of possible **ACTIONS** actions that will affect the **environment** and produce a new **state**.

A **REWARD**. A numerical value which reveals how positive or negative the **action** is.

You can't apply reinforcement learning if you're not able to define these objects properly!



# OBJECTIVE

Maximise the **long-term** rewards.

*Do not pull all the effort on capturing the queen if it means losing all you pieces.*

**How ?** There are two mains approaches:

**VALUE-BASED.** Look for the optimal reward.

- *Learn to estimate the expected rewards for each action in each state.*
- *Use this knowledge to choose the best action.*

**POLICY-BASED.** Look for the optimal **policy**.

- *Learn directly the best action to take for each observation.*

**NB:** Methods like Actor-Critic try to optimise both policy and rewards.

# OBJECTIVE

Maximise the **long-term** rewards.

*Do not pull all the effort on capturing the queen if it means losing all you pieces.*

**How ?** There are two mains approaches:

**VALUE-BASED.** Look for the optimal reward.

- *Learn to estimate the expected rewards for each action in each state.*
- *Use this knowledge to choose the best action.*

**POLICY-BASED.** Look for the optimal policy.

- *Learn directly the best action to take for each observation.*

**NB:** Methods like Actor-Critic try to optimise both policy and rewards.



# DEFINITIONS

---


# POLICY

The algorithm used by the agent to determine its action.

$$\Pi(s) = \operatorname{argmax}_a p(a/s)$$



It can be whatever you want :

**RULES.** *Example: Apply force to the right if the pole is tilting to left else apply force to the left.*



*Angle*  $< 0^\circ$   $\longrightarrow$  *Push to the right.*

## PARAMETRIC FUNCTION



*Push to the right.*

# POLICY SEARCH

How do you train the policy ?

## RANDOM SEARCH.

*Does not work when the space is too big, which is often the case.*

## GENETIC ALGORITHM

1. *Try a set of  $N$  policies.*
2. *Keep the  $n$  best policies*
3. *Generate  $N$  new policies that are random deviation of these  $n$  policies.*
4. *Iterate*

## POLICY GRADIENT

1. *Play the game with regards to the policy parameters.*
2. *Update these parameters with gradient descent.*


# THE CREDIT ASSIGNMENT PROBLEM

What is the difference with **(Deep) Q-Learning** algorithm?

Do not evaluate the **Q-value**.

Predict the **policy** and the **action to take directly**.

How to choose the loss on which to train the policy ?



**PROPOSITION:** the immediate reward ?

**PROBLEM:** we do not know the influence on **the long-term reward** of an action.

**EXAMPLE:** Going up is obviously not the best action.

**SOLUTION:** **DISCOUNTED CUMULATIVE EXPECTED REWARD**

# THE DISCOUNTED CUMULATIVE EXPECTED REWARD

Evaluate an action based on the sum of all the rewards that come after it.

$$R_t = \sum_{i=t}^T \gamma^i r_i$$

where  $\gamma$  is the discounted rate and  $r_t$  is the reward as step  $t$ .

**PacMan Example** *The agent decides to go up three times in a row. It gets +10 reward after the first step, 0 after the second step, and finally -50 after the third step (by being killed).*

With  $\gamma = 0.8$

Step	Immediate reward	Discounted cumulated expected reward
0	10	$R_0 = 10 + 0 \times 0.8 + (-50) \times 0.8^2 = -22$
1	0	$R_1 = 10 + (-50) \times 0.8 = -40$
2	-50	$R_2 = -50$

# POLICY GRADIENT ALGORITHM

---



# POLICY GRADIENT

**OBJECTIVE:** Find a **policy**  $\Pi_\theta$ , with parameters  $\theta$  that **maximises** the **expected** values of the sum of the discounted rewards.

$$J(\theta) = \mathbb{E}_{\pi_\theta} \left[ \sum_{t=0}^{T-1} \gamma^t r_t \right]$$

- The environment are usually **non-deterministic**.
- $\theta$  are the parameters of the neural network.
- Let's perform a **gradient ascent** search to learn the optimal  $\theta$ .

$$\theta \leftarrow \theta + \alpha \nabla J(\theta)$$

- $\alpha$  is the learning rate.

# GRADIENT SEARCH

How to find  $\nabla J(\theta)$  ? with:

$$\begin{aligned} J(\theta) &= \mathbb{E}_{\pi_{\theta}} \left[ \sum_{t=0}^{T-1} \gamma^t r_t \right] \\ &= \mathbb{E}_{\pi_{\theta}} [R(\tau)] \\ &= \sum_{\tau} P(\tau) R(\tau) \end{aligned}$$

Where is the **trajectory** of the agent moving through the environment

$$\tau = (s_0, a_0, r_0, s_1, a_1, r_1, \dots, s_{T-1}, a_{T-1}, r_{T-1}, s_T)$$

# LOG DERIVATIVE TRICK

Using the *log-derivative* trick:

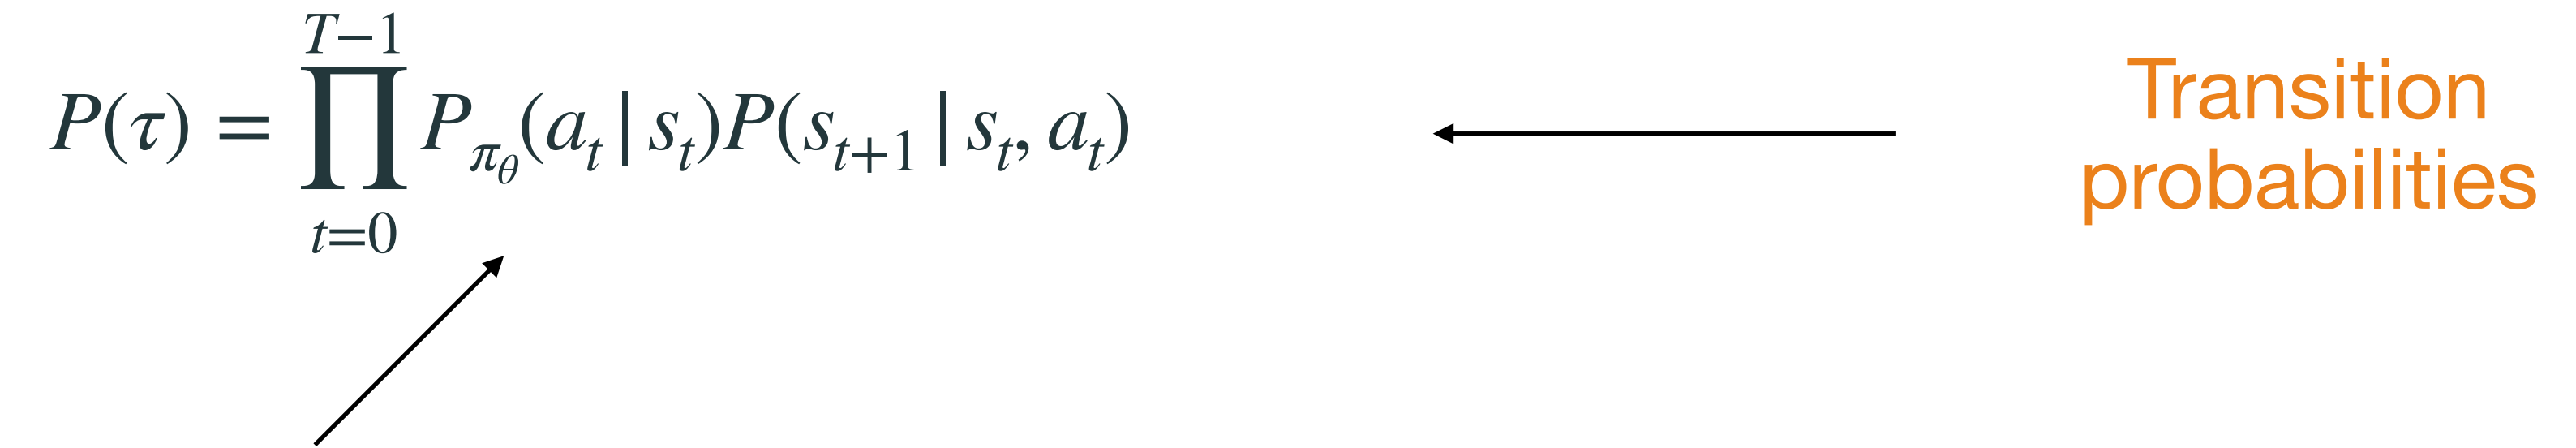
$$\begin{aligned}\nabla_{\theta} J(\theta) &= \sum_{\tau} \nabla_{\theta} P(\tau) R(\tau) \\ &= \sum_{\tau} P(\tau) \frac{\nabla_{\theta} P(\tau)}{P(\tau)} R(\tau) \\ &= \sum_{\tau} P(\tau) \nabla_{\theta} \log P(\tau) R(\tau) \\ &= \mathbb{E}[\nabla_{\theta} \log P(\tau) R(\tau)]\end{aligned}$$

During training, trajectories are *randomly sampling*.

To maximise the expectation above, we need to maximise it with respect to its argument i.e. we maximise:

$$\nabla_{\theta} J(\theta) \sim R(\tau) \nabla_{\theta} \log P(\tau)$$

# GRADIENT SEARCH

$$P(\tau) = \prod_{t=0}^{T-1} P_{\pi_{\theta}}(a_t | s_t) P(s_{t+1} | s_t, a_t)$$


Transition  
probabilities

Probability to choose an **action** at a given **state** according to the **policy** (Output of softmax)

$$\begin{aligned}\nabla_{\theta} \log P(\tau) &= \nabla \log \left( \prod_{t=0}^{T-1} P_{\pi_{\theta}}(a_t | s_t) P(s_{t+1} | s_t, a_t) \right) \\ &= \nabla_{\theta} \left[ \sum_{t=0}^{T-1} (\log P_{\pi_{\theta}}(a_t | s_t) + \log P(s_{t+1} | s_t, a_t)) \right] \\ &= \nabla_{\theta} \sum_{t=0}^{T-1} \log P_{\pi_{\theta}}(a_t | s_t)\end{aligned}$$

# GRADIENT SEARCH

$$\begin{aligned}\nabla_{\theta} J(\theta) &\sim R(\tau) \nabla_{\theta} \sum_{t=0}^{T-1} \log P(\tau) \\ &\sim R(\tau) \nabla_{\theta} \sum_{t=0}^{T-1} \log P_{\pi_{\theta}}(a_t | s_t)\end{aligned}$$

This is a **weighted cross entropy** where the weight are the discounted reward !

$$CE = - \sum p(x) \log(q(x))$$

Hence by training a neural network using **weighted cross entropy** with discounted reward as the weights

You are training a PG algorithm.

# KERAS - DISCOUNTED LOSS

```
import tensorflow.keras.losses as klo
class discountedLoss(klo.Loss):
    """
    Args:
        pos_weight: Scalar to affect the positive labels of the loss function.
        weight: Scalar to affect the entirety of the loss function.
        from_logits: Whether to compute loss from logits or the probability.
        reduction: Type of tf.keras.losses.Reduction to apply to loss.
        name: Name of the loss function.
    """

    def __init__(self,
                 reduction=klo.Reduction.AUTO,
                 name='discountedLoss'):
        super().__init__(reduction=reduction, name=name)

    def call(self, y_true, y_pred, adv):
        log_lik = - (y_true * K.log(y_pred) + (1 - y_true) * K.log(1 - y_pred))
        loss = K.mean(log_lik * adv, keepdims=True)
        return loss
```



# KERAS - CUSTOM MODEL CLASS

```
import tensorflow.keras.models as km
import tensorflow.keras.layers as kl
import tensorflow.keras.initializers as ki
import tensorflow.keras.metrics as kme

class kerasModel(km.Model):
    def __init__(self):
        super(kerasModel, self).__init__()
        self.layersList = []
        self.layersList.append(kl.Dense(9, activation="relu",
                                         input_shape=(4,),
                                         use_bias=False,
                                         kernel_initializer=ki.VarianceScaling(),
                                         name="dense_1"))
        self.layersList.append(kl.Dense(1,
                                         activation="sigmoid",
                                         kernel_initializer=ki.VarianceScaling(),
                                         use_bias=False,
                                         name="out"))

        self.loss = discountedLoss()
        self.optimizer = ko.Adam(lr=1e-2)
        self.train_loss = kme.Mean(name='train_loss')
        self.validation_loss = kme.Mean(name='val_loss')
        self.metric = kme.Accuracy(name="accuracy")

    @tf.function()
    def predict(x):
        """
        This is where we run
        through our whole dataset and return it, when training and testing.
        """
        for l in self.layersList:
            x = l(x)
        return x
    self.predict = predict

    @tf.function()
    def train_step(x, labels, adv):
        """
        This is a TensorFlow function, run once for each epoch for the
        whole input. We move forward first, then calculate gradients with
        Gradient Tape to move backwards.
        """
        with tf.GradientTape() as tape:
            predictions = self.predict(x)
            loss = self.loss.call(
                y_true=labels,
                y_pred = predictions,
                adv = adv)
            gradients = tape.gradient(loss, self.trainable_variables)
            self.optimizer.apply_gradients(zip(gradients, self.trainable_variables))
            self.train_loss(loss)
            return loss

        self.train_step = train_step
```

Define a **KERAS MODEL** class.

*optimiser*, *train\_loss*, *validation\_loss*, *metric* are native keras function.

*loss* is the custom loss we defined before.

*predict* and *train\_step* are redefine according to our needs

# REINFORCE ALGORITHM (R.WILLIAMS-1992)

1. Initiate the **policy** randomly.
2. Let the **policy** play the game several times. and at each step compute the **gradients** that would make the chosen action even, don't apply these **gradients** yet.
3. Compute each action's **discounted cumulative expected** reward
4. Multiply each **gradient** vector by the corresponding action's score.
5. Compute the mean of all the resulting **gradient** vectors, and use it to perform a Gradient Descent step.

# DIFFERENCE WITH Q-LEARNING

In (deep) Q-learning, the parameters we are trying to find are those that minimise the difference between the actual Q values (drawn from experiences) and the target.

No explicit exploration.

Converge faster, but need more and complete episode.

Still better with memory replay.

No target network needed

# EXPLORATION VS EXPLOITATION

**EXPLOITATION MODE** : Pick the **best action** according to the **policy**.

*Problem : Being stuck in a non-optimal solution.*

**EXPLORATION MODE** : Takes **random action** to explore the space.

*Problem : Can take a lot of time.*

## **E-GREEDY STRATEGY**

*Take the **best action**  $(1-\epsilon)\%$  of the time*

*Take a **random action**  $\epsilon\%$  of the time.*

**STOCHASTIC STRATEGY** : Use the probability to take an action to choose to act randomly or not.

*$action = 0$  if  $random.uniform(0, 1) < predict\_proba$ .*

*$action = 1$  otherwise.*

# PSEUDO CODE

- Initiate the model to train  $P$ .
- While  $num\_episode < max\_number\_episode$  OR  $test\_score < goal$ 
  - Play episode(s) **entirely** :
    - Choose action according to a strategy.
    - Save **experiences**.
  - If  $len(experiences) > batch\_size$ 
    - Normalise all **discounted rewards**
    - Train the model  $P$  with
      - $X = \text{States}$
      - $y = \text{actions}$
      - Loss = **discounted\_loss**

$Experiences = [state, action, discounted\_reward]$

Discounted reward are computed from **complete** episode (**Monte-Carlo** method).

# BASELINE

The **discounted rewards** are **normalised** before each batch.

$$\nabla_{\theta} J(\theta) \sim \mathbb{E}[(R(\tau) - N) \nabla_{\theta} \log P(\tau)]$$

We subtract a **baseline**. *It reduce the variance of gradient estimation while keeping the bias unchanged.*

A common **baseline** is to subtract **state-value**.

$$\nabla_{\theta} J(\theta) \sim \mathbb{E}[(R(\tau) - V(s)) \nabla_{\theta} \log P(\tau)]$$

And because  $Q^{\pi}(s, a) = \mathbb{E} \left( \lim_{H \rightarrow \infty} \sum_{t=0}^H \gamma^t r_t \middle| s_0 = s, a_0 = a, \pi \right)$

$$\nabla_{\theta} J(\theta) \sim \mathbb{E}[(Q(s, a) - V(s)) \nabla_{\theta} \log P(\tau)]$$




# ACTOR-CRITIC

- Initiate the **policy** model  $P_\theta$  and the **Q-value** model  $Q_\omega$  .
- While  $num\_episode < max\_number\_episode$  OR  $test\_score < goal$ 
  - Play a step of an episode :
    - Choose action according to a strategy (use P in exploit mode).
    - Save experiences
  - Train the **policy** model
    - $\theta \leftarrow \theta + \alpha \nabla Q(s, a) \log P_{\pi_\theta}(a_t | s_t)$
  - Train the **Q-value** model
    - Compute target from experiences
    - $\omega \leftarrow \omega - \alpha \nabla_\omega \mathbb{E}_{s' \sim P(s'|s,a)} [(Q_\omega(s, a) - target(s'))^2]$

One Notebook : *Policy Gradient.ipynb*

- Implement hard coded policy
- Train a neural network model to learn a given policy.
- Train a neural model with a **Policy Gradient algorithm**.
  - Implement **discounted reward** function.
  - Implement **loss** using **keras**.
  - Write PG algorithm.



<https://adventuresinmachinelearning.com/policy-gradient-tensorflow-2/>

[https://medium.com/@thechrisyoon/deriving-policy-gradients-and-implementing-reinforce-f887949bd63#:~:text=REINFORCE%20is%20a%20Monte%2DCarlo,to%20update%20the%20policy%20parameter.&text=Store%20log%20probabilities%20\(of%20policy\)%20and%20reward%20values%20at%20each%20step](https://medium.com/@thechrisyoon/deriving-policy-gradients-and-implementing-reinforce-f887949bd63#:~:text=REINFORCE%20is%20a%20Monte%2DCarlo,to%20update%20the%20policy%20parameter.&text=Store%20log%20probabilities%20(of%20policy)%20and%20reward%20values%20at%20each%20step)

<https://towardsdatascience.com/understanding-actor-critic-methods-931b97b6df3f>

<https://lilianweng.github.io/lil-log/2018/04/08/policy-gradient-algorithms.html>

# REFERENCES

Géron, A. (2017). Hands-On Machine Learning with Scikit-Learn and TensorFlow: Concepts. *Tools, and Techniques to build intelligent systems*.

Williams, R. J. (1992). Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine learning*, 8(3-4), 229-256.